

PA: vector Abstract Data Type

Learning Outcomes

- Gain experience in data abstraction and encapsulation techniques using classes.
- Gain experience in encapsulating pointers, dynamic arrays, and free store in classes.
- Gain practice in debugging memory leaks and errors using Valgrind.
- Read and understand code.

Task

The purpose of this laboratory is to understand how to build our own version of the most useful container in the standard library: `std::vector`. A `std::vector` provides a sequence of elements of a given type; you can refer to an element by its index; extend the `std::vector` by using `push_back`; ask a `vector` for the number of its elements using `size`; and have access to the `std::vector` checked against attempts to access out-of-range element. `std::vector` is a convenient, flexible, and an efficient [in both time and space] container of elements. `std::string` has similar properties, as do other useful standard container types, such as `std::list` and `std::map`. Learning how to build our own version of `std::vector` from the basic language facilities available to every programmer allows us to illustrate useful concepts and programming techniques. The techniques we encounter in the implementation of a `std::vector` clone are generally useful and very widely used.

In many ways, the `str` class we've implemented in this course has similar properties as a `std::vector<Str>` is a container of `char` s while our `vector` is a container of `int` s. Using the knowledge gained in the implementation of `Str` class [source code is available on course web page] and using the following partial class in `vector.hpp`

```

1  #include <cstddef>           // for size_t
2  #include <initializer_list> // for std::initializer_list<int>
3
4  namespace hlp2 {
5
6  class vector {
7  public:
8      // type definitions for value_type, size_type,
9      // pointer, const_pointer, reference, const_reference
10 public:
11     vector(); // default ctor
12     explicit vector(size_type n); // non-default ctor
13     vector(std::initializer_list<int> rhs); // non-default ctor
14     vector(vector const& rhs); // copy ctor
15     ~vector();
16
17     vector& operator=(vector const&);
18     vector& operator=(std::initializer_list<int> rhs);
19
20     reference operator[](size_type index);
21     const_reference operator[](size_type index) const;
22
23     void reserve(size_type n);

```

```

24 void resize(size_type n);
25 void push_back(value_type val);
26
27 bool empty() const; // is container empty?
28 size_type size() const; // what is sz?
29 size_type capacity() const; // what is space?
30 size_type allocations() const; // how many allocations or "growths"?
31
32 // iterators ...
33 pointer begin(); // pointer to first element
34 pointer end(); // pointer to one past last element
35 const_pointer begin() const;
36 const_pointer end() const;
37 const_pointer cbegin() const;
38 const_pointer cend() const;
39
40 private:
41     size_type sz; // number of elements
42     size_type space; // number of elements plus "free space"/"slots"
43     size_type allocs; // number of "growths"
44     pointer data; // address of first element
45 };
46
47 }

```

you must define the member functions in `vector.cpp` such that these functions successfully pass the unit tests implemented in the driver. If you know how to define the class functions, you can proceed to implementing the functions. If you want to get a better understanding, an explanation is provided [here](#). In either case, heed this warning:

Be particularly careful about managing memory to avoid memory leaks and errors. Since there are multiple objects in memory, allocation and deallocation is a complex process. If the TA finds a memory leak or error in your submission, you will receive a maximum grade of D.

Submission Details

Please read the following details carefully and adhere to all requirements to avoid unnecessary deductions.

Submission files

You will be submitting interface file `vector.hpp`. Complete the definition of class `vector` and upload to the submission server. This file should not include `<vector>` nor contain the string `<vector>`.

Complete the necessary definitions in `vector.cpp` and upload to the submission server. This file should not include `<vector>` nor contain the string `<vector>`.

Compiling, executing, and testing

Download `vector-driver.cpp` and a correct input file `output.txt` for the unit tests in the driver. Follow the steps from the previous lab to refactor a `makefile` and test your program.

Remember, to frequently check your code for memory leaks and errors with Valgrind. Also note that if your submission has even a single memory leak or error, D will be the maximum possible grade. You're warned!!!

File-level and function-level documentation

Every source and header file *must* begin with a *file-level* documentation block. This module will use [Doxygen](#) to tag source and header files for generating html-based documentation. In addition, every function that you declare and define and submit for assessment must contain *function-level documentation*. This documentation should consist of a description of the function, the inputs, and return value.

Submission and automatic evaluation

1. In the course web page, click on the appropriate submission page to submit the necessary files.
2. Please read the following rubrics to maximize your grade. Your submission will receive:
 - F grade if your submission doesn't compile with the full suite of `g++` options.
 - F grade if your submission doesn't link to create an executable.
 - Your implementation's output must exactly match correct output of the grader (you can see the inputs and outputs of the auto grader's tests). There are only two grades possible: $A+$ grade if your output matches correct output of auto grader; otherwise F .
 - A maximum of D grade if Valgrind detects even a single memory leak or error. A teaching assistance will check you submission for such errors.
 - A deduction of one letter grade for each missing documentation block in your submissions. Each source file must have **one** file-level documentation block and a function-level documentation block for each defined function. A teaching assistant will physically read submitted source files to ensure that these documentation blocks are authored correctly. Each missing or incomplete or copy-pasted (with irrelevant information from some previous assessment) block will result in a deduction of a letter grade. For example, if the automatic grader gave your submission an $A+$ grade and one documentation block is missing, your grade will be later reduced from $A+$ to $B+$. Another example: if the automatic grade gave your submission a C grade and the two documentation blocks are missing, your grade will be later reduced from C to F .

How to "grow" vector?

Essential functions

The most useful container in the C++ standard library is `vector`. A `vector` provides a sequence of elements of a given type. You can refer to an element by its index, extend the `vector` by using `push_back`, ask a `vector` for the number of its elements using `size`, and have access to the `vector` checked against attempts to access out-of-range element. The standard library `vector` is a convenient, flexible, and an efficient (in both time and space) container of elements. The standard `string` has similar properties, as have other useful standard container types, such as `list` and `map`. Learning how to build our own version of `vector` from the basic language facilities available to every programmer allows us to illustrate useful concepts and programming techniques. The techniques we encounter in the `vector` implementation are generally useful and very widely used.

The `string` class we've implemented in this course has similar properties as a `vector`. Using the following partial class in `vector.h`

```

1  #include <cstddef> // for size_t
2
3  namespace hlp2 {
4  class vector {
5  public:
6      using size_type = size_t;
7      using value_type = int;
8      using pointer    = value_type*;
9      // other type definitions for const_pointer,
10     // reference, const_reference, ...
11  public:
12     // interface
13  private:
14     size_type sz; // number of elements
15     pointer data; // pointer to 1st element of free store memory block
16 };
17 }

```

you should be able to (easily) implement the following essential functions in `vector.cpp`:

1. Default constructor: `vector::vector();`
2. Non-default constructor that allocates memory for `n` elements with each element initialized to 0: `explicit vector::vector(int n);`
3. Copy constructor: `vector::vector(vector const& rhs);`
4. Copy assignment (copy object of same type): `vector& vector::operator=(vector const& rhs);`
5. Destructor: `vector::~~vector();`
6. Subscript operator: `reference vector::operator[](size_type index);`
7. Subscript operator overloaded on `const`: `const_reference vector::operator[](size_type index) const;`
8. Accessor function to return the number of `int` elements encapsulated by the container: `size_type vector::size() const;`

At this point, we can:

- Create `vector`s of `int`s with whatever number of elements we want
- Copy our `vector`s using assignment and initialization
- Rely on `vector`s to correctly release their memory when they go out of scope
- Access `vector` elements using the conventional subscript notation on both the right-hand and left-hand sides of an assignment
- Iterate through `vector` elements using its size

Growing a `vector`

That's all good and useful, but to reach the level of sophistication we expect (based on our experience with `std::vector`), we need to address three more concerns:

- How do we change the size of a `vector`?
- How do we specify the element type of a `vector` as an argument?
- How do we catch and report out-of-range `vector` element access similar to the `at` member function of `std::vector`?

Specifying the element type as an argument is related to the topic of templates while reporting out-of-range element access is related to the topic of exceptions. These topics will be discussed in the second half of the course. This lab is concerned with changing a `vector`'s size.

The standard library `std::vector` offers a variety of operations that can change a `vector`'s size, such as `push_back`, `resize`, assignment `=`, `insert`, `erase`, `pop_back`, and so on. Consider the following code fragment involving the standard library `std::vector` that reads an unknown number of `int` values:

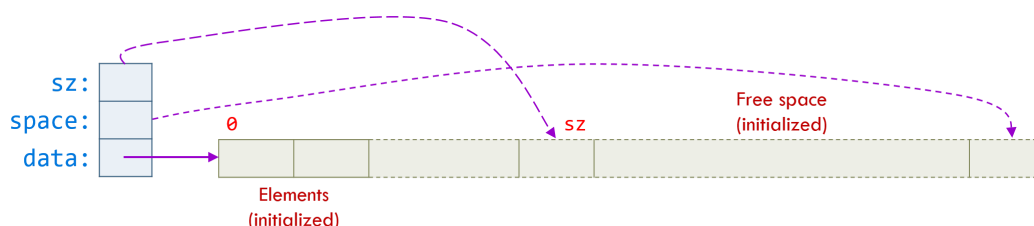
```
1 std::vector<int> vi;
2 for (int i; std::cin >> i; ) {
3     vi.push_back(i);
4 }
```

Container `vi` begins its life with zero size. Here's one possibility for the unknown number of `int` values to be stored in the container. In the first iteration, `push_back` will allocate sufficient memory to store the first `int` value. In the second iteration, `push_back` finds no room for the second element and therefore has to allocate a new block of memory to store both elements. This is a simple strategy for changing size: just allocate memory for the new number of elements, copy old elements into the new memory, and then delete the old memory.

If you resize often, that's inefficient. So, we can optimize our programs by anticipating such changes in size by keeping track of both the number of elements and the amount of "free space" reserved for "future expansion." The `vector` class would then look like this:

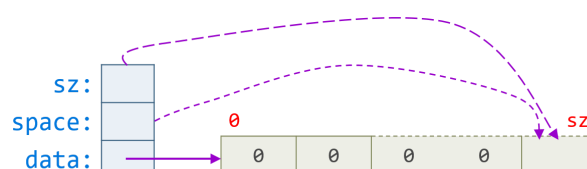
```
1 namespace hlp2 {
2     class vector {
3         size_type sz;    // number of elements
4         size_type space; // number of elements plus "free space"/"slots"
5         pointer data;    // address of first element
6     public:
7         // interface here ...
8     };
9 } // end namespace hlp2
```

We can represent this graphically like this:



Since elements are indexed from `0`, we represent `sz` (the number of elements) as referring to one beyond the last element and `space` as referring to one beyond the last allocated slot. The pointers shown are really `data+sz` and `data+space`.

When a `vector` is constructed with non-default constructor, `sz==space`; that is, there is no "free space":



We don't start allocating extra slots until we begin changing the number of elements using member function `push_back`.

When a `vector` is default constructed like this:

```
1 | hlp2::vector vi;
```


`sz` and `space` are both 0 and `data` is `nullptr`:

<code>sz:</code>	0
<code>space:</code>	0
<code>data:</code>	<code>nullptr</code>

Now, when we use member function `push_back` like this:

```
1 | vi.push_back(5);
```

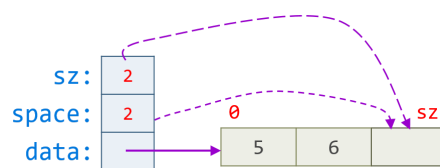
`push_back` must "grow" the empty container by one element:



Calling `push_back` again:

```
1 | vi.push_back(6);
```

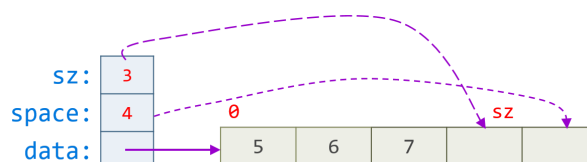
will cause `push_back` to "grow" the container since `sz==space==1`. This time, the function will "grow" the container by twice its previous size to two elements. This necessitates allocating a block of memory for two `int`s; copying the old element to index 0 of the new block; adding the new element with value 6 to the end of the new block; followed by deleting the old block of memory. After the conclusion of the "growing" process, the container looks like this:



Calling `push_back` a third time:

```
1 | vi.push_back(7);
```

will require the function to "grow" the container again since `sz==space==2`. The function will again "grow" the container by twice its previous size - from two to four elements. In the memory block, three elements are used with the fourth slot available for "future growth." This is illustrated in the following picture:



reserve and capacity

The most fundamental function when we change sizes (that is, when we change the number of elements) is `vector::reserve`. This is the function that will be used by both clients and other functions in the `vector` interface to add memory for new elements:

```

1  // Requests that capacity be at least enough to contain n elements.
2  // If n is greater than current capacity, function causes container
3  // to reallocate its storage increasing its capacity to n (or greater).
4  // In all other cases, the function call does not cause a reallocation
5  // and the vector capacity is not affected.
6  void vector::reserve(vector::size_type newalloc) {
7      // If newalloc <= space, return because the function will never
8      // decrease allocation
9      // Allocate new block of memory to store newalloc number of ints
10     // Note: we don't initialize elements of reserved space. After all,
11     // we're just reserving space; using that space for elements is the
12     // job of functions such as push_back and resize
13     // Copy old elements
14     // Delete memory containing old elements pointed to by data
15     // Assign to data address of first int element in new memory block
16     // Set space to newalloc: space = newalloc;
17     // increment allocs (by one)
18 }

```

This tutorial will measure the number of memory allocations in your submissions. To keep track of the count of memory allocations, the `vector` class is augmented with a data member `allocs` that is initialized to zero during construction and then incremented each time function `reserve` is called:

```

1  namespace hlp2 {
2  class vector {
3      size_type sz;    // number of elements
4      size_type space; // number of elements plus "free space"/"slots"
5      size_type allocs; // memory allocation counter
6      pointer    data; // address of first element
7  public:
8      // interface here ...
9  };
10 } // end namespace hlp2

```

Obviously, the amount of "free space" available for "future growth" is of interest to a client, so we (just like the standard) provide an accessor member function `size_type capacity() const;` for obtaining that information. That is, for a `vector` called `vi`, `vi.capacity()-vi.size()` is the number of elements we could `push_back` to `vi` without causing reallocation.

resize

Member function `void vector::resize(size_type n);` resizes the container so that it contains `n` elements. We've to handle several cases:

1. The new size `n` is larger than the old allocation `space`. In this case, an automatic reallocation takes place to "grow" the container so that both `sz` and `space` are `n`. The new elements are value-initialized which for `int`s means the elements have `0` value.

2. The new size `n` is larger than the old size `sz`, but smaller than or equal to old allocation `space`. In this case, the "free space" is used to accommodate the new elements that are value-initialized (in the case of `int`s to `0`). Size `sz` changes to `n`.
3. The new size `n` is equal to the old size `sz`. Nothing happens.
4. The new size `n` is smaller than the old size `sz`. The number of elements is reduced to the first `n` elements. Unfortunately, unlike the standard library version, we're unable to destroy the elements. The intricacies of memory allocation/deallocation and destruction is beyond the scope of this course. Instead, we leave the elements alone and just reduce size `sz` to `n`.

This function requires no more than 3 lines of code with member function `vector::reserve` doing the hard work of dealing with memory (for the first case above). The second case is easily dealt with by initializing each new element with the default value `0`. The third case doesn't require any work. The fourth case simply reduces the container's size to `n`.

push_back

When we first think of it, member function `push_back` may appear complicated to implement, but given `reserve`, it is quite simple:

```

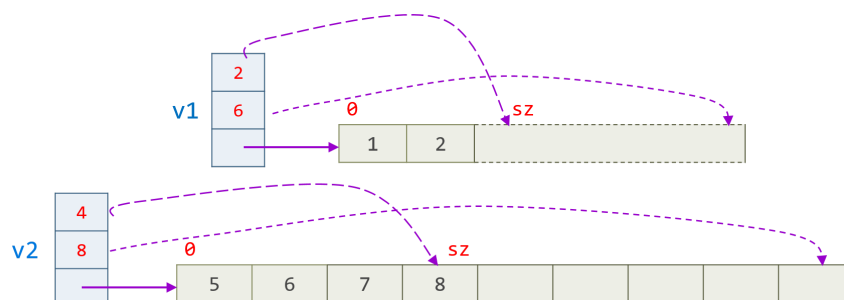
1 // Adds a new element at the end of the vector, after its current last
2 // element. The content of val is copied (or moved) to the new element.
3 // This effectively increases the container size by one, which causes an
4 // automatic reallocation of allocated storage space if - and only if -
5 // the new vector size surpasses the current vector capacity.
6 void vector::push_back(value_type val) {
7     // if space==0 then reserve space for 1 element
8     // else if sz==space then reserve space for 2*space elements
9     // add val to data[sz]
10    // increment sz
11 }

```

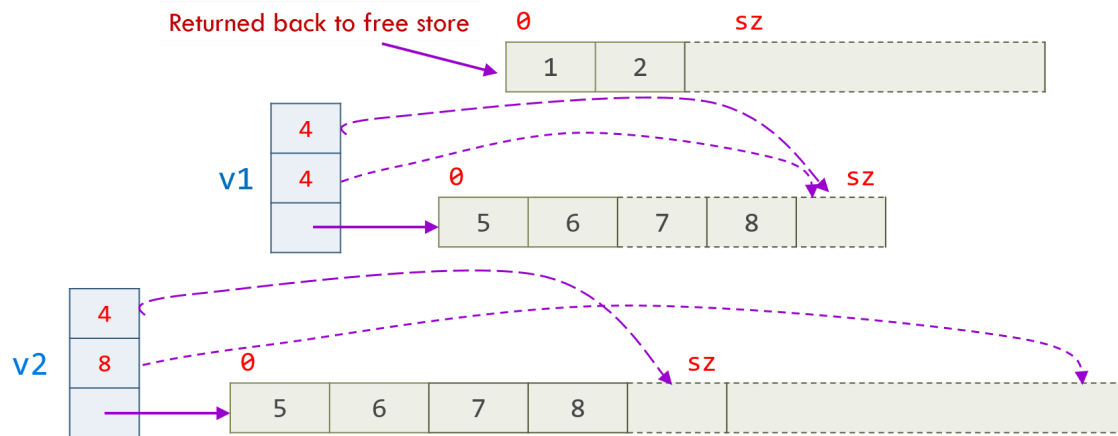
In other words, if we've no spare space, we double the size of the allocation. In practice, this turns out to be a very good choice for the vast majority of uses of `vector`, and that's the strategy used by most implementation of the standard library `std::vector`.

Copy assignment operator

We know how to write copy assignments. Obviously, we need to copy the elements, but what about the spare memory slots? Do we "copy" the "free space" at the end? Consider:



We don't copy the "free space" at the end: the new `vector` will get a copy of the elements, but since we have no idea how that new `vector` is going to be used, we don't bother with extra space at the end:



The simplest implementation of that is:

- Allocate memory for a copy
- Copy the elements
- Delete the old allocation
- Set `data` to the new value
- Set `sz` and `space` to new values

This implementation is correct, but when we look at it a bit we realize that we do a lot of redundant allocation and deallocation. What if the `vector` we assign to has more elements than the one we assign? What if the `vector` we assign to has the same number of elements as the `vector` we assign? In many applications, that last case is very common. In either case, we can just copy the elements into the memory already available in the target `vector`. This optimization is left as an exercise.

Caution: The unit tests use the simplest implementation described above and do not take into account any of the optimizations described here.